



Alzheimer's Disease Diagnosis with Deep Learning

Machine Learning Project

The Bridge Data Science Bootcamp

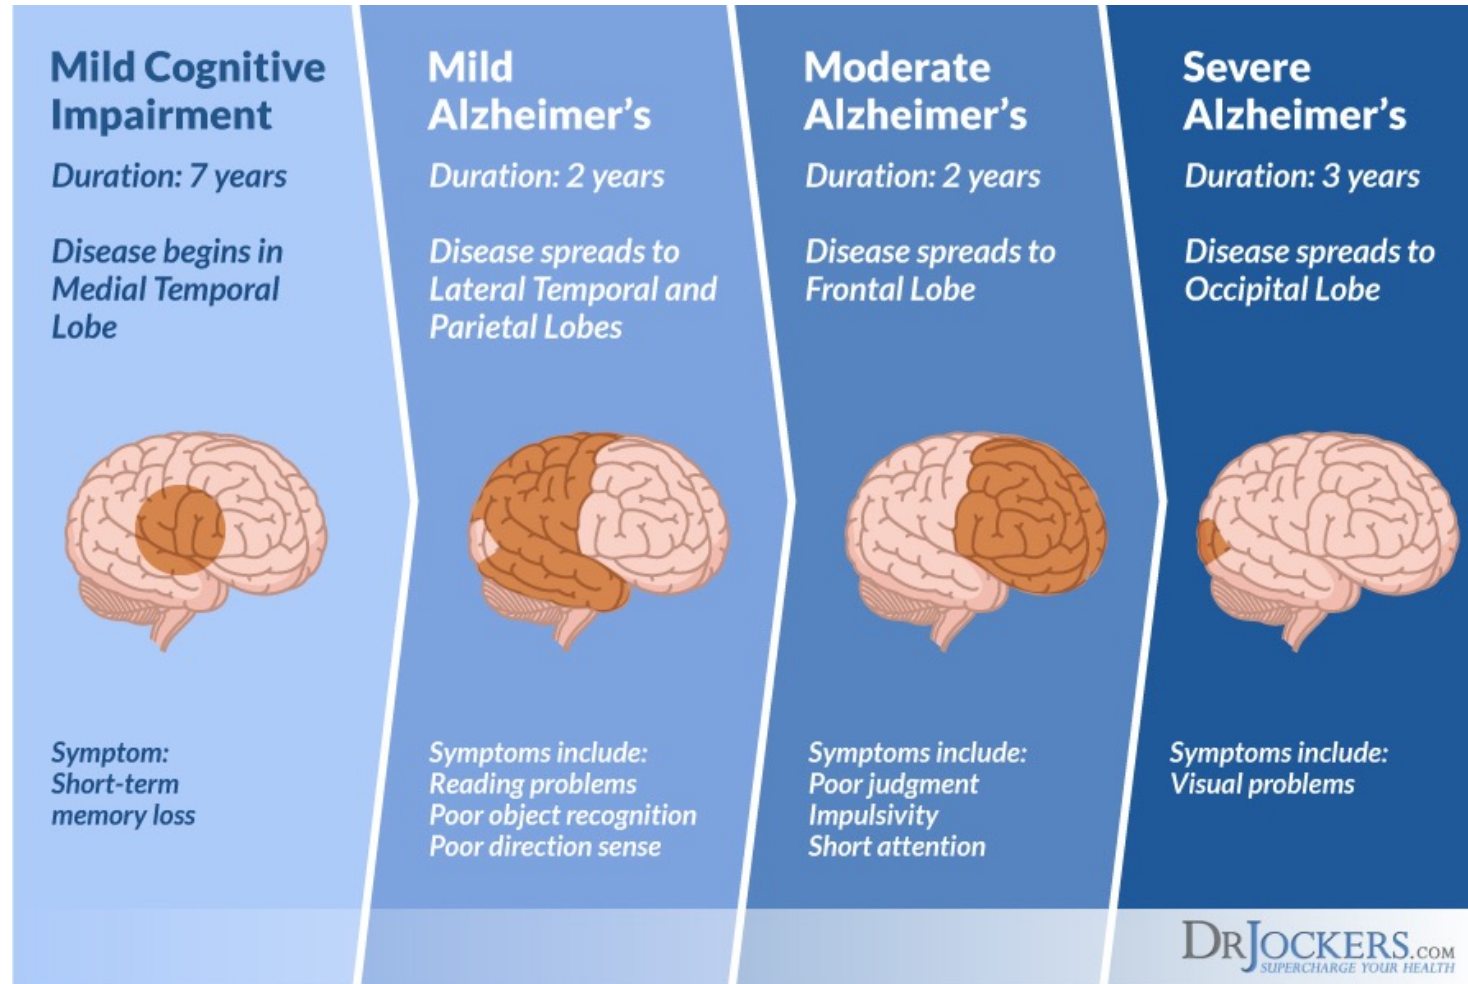
Sara Picó

09.04.2024

Problem statement



Different stages




Early diagnosis is crucial

How?




MAGNETIC RESONANCE IMAGING



Objective


Develop and evaluate a **Deep Learning** model capable of **identifying distinctive patterns** in MRI images to discriminate between different stages




Content

Training and test sets (4 classes):

- Mild Demented
- Moderate Demented
- Non Demented
- Very Mild Demented





Methodology

1. Data preparation
2. Model comparison
3. Tune the model
4. Evaluate performance
5. Final conclusions

1. Data Preparation

Import libraries

Built helper functions:

- read_data
- show_random_images_with_labels
- flatten_gray_scale
- plot_training_metrics
- prepare_for_test
- show_images




1. Data Preparation


View sample images:

```
show_random_images_with_labels(X_train, y_train, num_images=10)
```


Label: NonDemented




Label: NonDemented




Label: NonDemented




Label: VeryMildDemented




Label: NonDemented




Label: VeryMildDemented




Label: VeryMildDemented




Label: MildDemented



Label: NonDemented



Label: VeryMildDemented



1. Data Preparation

View proportion of classes:



Warning! Is unbalance

We should check AUC instead of only Accuracy

2. Model comparison

Random Forest


Basic Convolutional model

Transfer Learning:

- ResNet50
- InceptionV3
- DenseNet169
- VGG16



2. Model comparison



Random Forest


- Flatten X_train
- Apply shuffle
- Apply model

```
rf_clf = RandomForestClassifier()  
np.mean(cross_val_score(rf_clf, X_train_rf_shuffled, y_train_shuffled,  
cv = 5, scoring = "accuracy"))
```

92% accuracy



2. Model comparison



Simple convolutional model

```
datagen = IDG(rescale = 1./255, validation_split=0.1)

train_gen = datagen.flow_from_directory(directory=train_dir,
                                        target_size=DIM,
                                        batch_size=400,
                                        class_mode='categorical',
                                        subset='training',
                                        shuffle=True)


validation_gen = datagen.flow_from_directory(directory=train_dir,
                                              target_size=DIM,
                                              batch_size=400,
                                              class_mode='categorical',
                                              subset='validation',
                                              shuffle=True)

test_gen = datagen.flow_from_directory(directory=test_dir,
                                       target_size=DIM,
                                       batch_size=6400,
                                       class_mode='categorical')
```

```
Found 4610 images belonging to 4 classes.
Found 511 images belonging to 4 classes.
Found 1279 images belonging to 4 classes.
```



2. Model comparison



Simple convolutional model

```
# Create the model
model = Sequential()

# Add convolutional layer
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_SIZE, IMG_SIZE, 3)))
# Add pooling layer
model.add(MaxPooling2D((2, 2)))

# Add another convolutional and pooling layers
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D((2, 2)))

# Flatten for the dense layer
model.add(Flatten())

# Add dense layer
model.add(Dense(128, activation='relu'))


# Out layer with softmax activation for classification
model.add(Dense(len(CLASSES), activation='softmax'))

# Compile the model with personalized metrics
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=[AUC(name='auc'), Precision(name='precision'), Recall(name='recall')])


# Train the model
history = model.fit(train_gen,
                    epochs=50,
                    validation_data=validation_gen)
```





2. Model comparison



Simple convolutional model



Model AUC 79.03%, Accuracy 58.87% on Test Data				
	precision	recall	f1-score	support
NonDemented	0.13	0.13	0.13	71
VeryMildDemented	0.00	0.00	0.00	5
MildDemented	0.47	0.46	0.47	256
ModerateDemented	0.32	0.34	0.33	179
accuracy			0.37	511
macro avg	0.23	0.23	0.23	511
weighted avg	0.37	0.37	0.37	511



Confusion Matrix
AUC: 79.03%

True label	Predicted label			
	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	9	1	40	21
VeryMildDemented	1	0	3	1
MildDemented	32	1	118	105
ModerateDemented	26	2	90	61

2. Model comparison

Transfer Learning: ResNet50

```
rn = ResNet50(input_shape=(IMG_SIZE,IMG_SIZE,3), weights='imagenet', include_top=False)
for layer in rn.layers:
    layer.trainable = False
x = Flatten()(rn.output)

prediction = Dense(4, activation='softmax')(x)

model = Model(inputs=rn.input, outputs=prediction)

model.compile(optimizer='adam',
              loss=tensorflow.losses.CategoricalCrossentropy(),
              metrics=[keras.metrics.AUC(name='auc'), 'acc'])
callback = keras.callbacks.EarlyStopping(monitor='val_loss',
                                         patience=8,
                                         restore_best_weights=True)

tic = time.perf_counter()
history = model.fit(train_gen,
                   steps_per_epoch=len(train_gen),
                   validation_data=validation_gen,
                   validation_steps=len(validation_gen),
                   epochs=50, callbacks=callback)
# time
toc = time.perf_counter()
print("Total Time:{}".format(round((toc-tic)/60,2)))
```



2. Model comparison


Transfer Learning: ResNet50



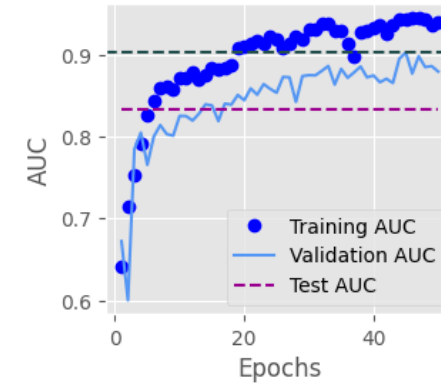
Model AUC 83.29%, Accuracy 53.95% on Test Data

	precision	recall	f1-score	support
NonDemented	0.40	0.17	0.24	179
VeryMildDemented	1.00	0.08	0.15	12
MildDemented	0.83	0.45	0.58	640
ModerateDemented	0.43	0.83	0.57	448
accuracy			0.54	1279
macro avg	0.66	0.38	0.39	1279
weighted avg	0.63	0.54	0.53	1279

Validation Loss by Epochs



Validation AUC by Epochs



Confusion Matrix
AUC: 83.29%

True label \ Predicted label	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	31	0	10	138
VeryMildDemented	1	1	0	10
MildDemented	18	0	288	334
ModerateDemented	28	0	50	370

2. Model comparison

Transfer Learning: InceptionV3

```
inception = InceptionV3(input_shape=(IMG_SIZE,IMG_SIZE,3), weights='imagenet', include_top=False)
for layer in inception.layers:
    layer.trainable = False
x = Flatten()(inception.output)

prediction = Dense(4, activation='softmax')(x)

model = Model(inputs=inception.input, outputs=prediction)

model.compile(optimizer='adam',
              loss=tensorflow.losses.CategoricalCrossentropy(),
              metrics=[keras.metrics.AUC(name='auc'), 'acc'])
callback = keras.callbacks.EarlyStopping(monitor='val_loss',
                                         patience=8,
                                         restore_best_weights=True)

tic = time.perf_counter()
history = model.fit(train_gen,
                   steps_per_epoch=len(train_gen),
                   validation_data=validation_gen,
                   validation_steps=len(validation_gen),
                   epochs=50, callbacks=callback)

# time
toc = time.perf_counter()
print("Total Time:{}".format(round((toc-tic)/60,2)))
```



2. Model comparison


Transfer Learning: InceptionV3




Model AUC 84.00%, Accuracy 60.20% on Test Data

	precision	recall	f1-score	support
NonDemented	0.73	0.20	0.31	179
VeryMildDemented	1.00	0.17	0.29	12
MildDemented	0.81	0.54	0.65	640
ModerateDemented	0.48	0.86	0.62	448
accuracy			0.60	1279
macro avg	0.76	0.44	0.47	1279
weighted avg	0.69	0.60	0.59	1279

Validation Loss by Epochs



Validation AUC by Epochs



Confusion Matrix
AUC: 84.00%

True label \ Predicted label	Predicted label			
	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	35	0	26	118
VeryMildDemented	0	2	3	7
MildDemented	4	0	347	289
ModerateDemented	9	0	53	386

2. Model comparison

Transfer Learning: DenseNet169

```
dense = DenseNet169(input_shape=(IMG_SIZE,IMG_SIZE,3), weights='imagenet', include_top=False)
for layer in dense.layers:
    layer.trainable = False
x = Flatten()(dense.output)

prediction = Dense(4, activation='softmax')(x)

model = Model(inputs=dense.input, outputs=prediction)

model.compile(optimizer='adam',
              loss=tensorflow.losses.CategoricalCrossentropy(),
              metrics=[keras.metrics.AUC(name='auc'), 'acc'])
callback = keras.callbacks.EarlyStopping(monitor='val_loss',
                                         patience=8,
                                         restore_best_weights=True)


tic = time.perf_counter()
history = model.fit(train_gen,
                   steps_per_epoch=len(train_gen),
                   validation_data=validation_gen,
                   validation_steps=len(validation_gen),
                   epochs=50, callbacks=callback)

# time
toc = time.perf_counter()
print("Total Time:{}".format(round((toc-tic)/60,2)))
```





2. Model comparison

Transfer Learning: DenseNet169



Model AUC 86.79%, Accuracy 62.39% on Test Data

	precision	recall	f1-score	support
NonDemented	0.92	0.06	0.12	179
VeryMildDemented	1.00	0.08	0.15	12
MildDemented	0.61	0.95	0.74	640
ModerateDemented	0.66	0.40	0.50	448
accuracy			0.62	1279
macro avg	0.80	0.37	0.38	1279
weighted avg	0.67	0.62	0.56	1279



Confusion Matrix
AUC: 86.79%

	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	11	0	117	51
VeryMildDemented	0	1	2	9
MildDemented	0	0	608	32
ModerateDemented	1	0	269	178

True label

Predicted label

2. Model comparison

Transfer Learning: VGG16

```
vgg = VGG16(input_shape=(IMG_SIZE,IMG_SIZE,3), weights='imagenet', include_top=False)
for layer in vgg.layers:
    layer.trainable = False
x = Flatten()(vgg.output)

prediction = Dense(4, activation='softmax')(x)

model = Model(inputs=vgg.input, outputs=prediction)
model.summary()

model.compile(optimizer='adam',
              loss=tensorflow.losses.CategoricalCrossentropy(),
              metrics=[keras.metrics.AUC(name='auc'), 'acc'])
callback = keras.callbacks.EarlyStopping(monitor='val_loss',
                                         patience=3,
                                         restore_best_weights=True)

tic = time.perf_counter()
history = model.fit(train_gen,
                   steps_per_epoch=len(train_gen),
                   validation_data=validation_gen,
                   validation_steps=len(validation_gen),
                   epochs=20, callbacks=callback)

# time
toc = time.perf_counter()
print("Total Time:{}".format(round((toc-tic)/60,2)))
```



2. Model comparison

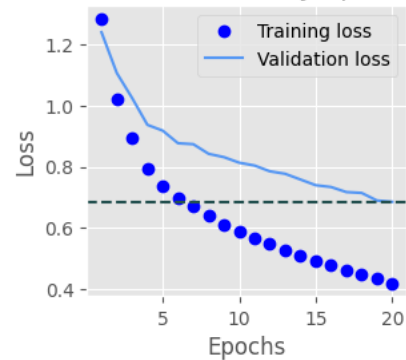
Transfer Learning: VGG16




Model AUC 89.00%, Accuracy 65.91% on Test Data

	precision	recall	f1-score	support
NonDemented	0.83	0.11	0.20	179
VeryMildDemented	1.00	0.17	0.29	12
MildDemented	0.67	0.88	0.76	640
ModerateDemented	0.63	0.57	0.60	448
accuracy			0.66	1279
macro avg	0.78	0.43	0.46	1279
weighted avg	0.68	0.66	0.62	1279

Validation Loss by Epochs



Validation AUC by Epochs



Confusion Matrix
AUC: 89.00%

True label	Predicted label			
	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	20	0	88	71
VeryMildDemented	0	2	2	8
MildDemented	1	0	566	73
ModerateDemented	3	0	190	255



3. Tune the model

DenseNet169 Fine-tuned

```
# Unfreeze the last few convolutional blocks for fine-tuning
for layer in dense.layers[:10]: # Unfreeze the last 10 layers for fine-tuning
    layer.trainable = True
```

Model AUC 87.87%, Accuracy 71.38% on Test Data

	precision	recall	f1-score	support
NonDemented	0.82	0.39	0.52	179
VeryMildDemented	0.36	0.42	0.38	12
MildDemented	0.80	0.75	0.77	640
ModerateDemented	0.62	0.81	0.70	448
accuracy			0.71	1279
macro avg	0.65	0.59	0.60	1279
weighted avg	0.74	0.71	0.71	1279



Confusion Matrix
AUC: 87.87%

True label \ Predicted label	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	69	2	44	64
VeryMildDemented	2	5	0	5
MildDemented	2	5	477	156
ModerateDemented	11	2	73	362



3. Tune the model

VGG16 Fine-tuned

```
# Freeze all layers except the last convolutional block
for layer in vgg.layers[:4]:
    layer.trainable = False
```

Model AUC 86.99%, Accuracy 62.47% on Test Data

	precision	recall	f1-score	support
NonDemented	0.56	0.21	0.31	179
VeryMildDemented	0.50	0.08	0.14	12
MildDemented	0.76	0.65	0.70	640
ModerateDemented	0.52	0.76	0.62	448
accuracy			0.62	1279
macro avg	0.58	0.43	0.44	1279
weighted avg	0.64	0.62	0.61	1279



Confusion Matrix
AUC: 86.99%

	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	38	0	45	96
VeryMildDemented	0	1	0	11
MildDemented	12	0	419	209
ModerateDemented	18	1	88	341

True label

Predicted label

3. Tune the model

DenseNet169 Fine-tuned


```
# Unfreeze the last few convolutional blocks for fine-tuning
for layer in dense.layers[:10]: # Unfreeze the last 10 layers for fine-tuning
    layer.trainable = True
```

Model AUC 87.87%, Accuracy 71.38% on Test Data

	precision	recall	f1-score	support
NonDemented	0.82	0.39	0.52	
VeryMildDemented	0.36	0.42	0.38	
MildDemented	0.80	0.75	0.77	
ModerateDemented	0.62	0.81	0.70	
accuracy			0.71	
macro avg	0.65	0.59	0.60	
weighted avg	0.74	0.71	0.71	



Training AUC by Epochs



Confusion Matrix
AUC: 87.87%


True label \ Predicted label				
	NonDemented	VeryMildDemented	MildDemented	ModerateDemented
NonDemented	69	2	44	64
VeryMildDemented	2	5	0	5
MildDemented	2	5	477	156
ModerateDemented	11	2	73	362

4. Evaluate with test images


DenseNet169
Fine-tuned

`show_images(train_gen, y_pred)`


Actual:ModerateDemented
Predicted:ModerateDemented




Actual:ModerateDemented
Predicted:ModerateDemented




Actual:ModerateDemented
Predicted:MildDemented




Actual:MildDemented
Predicted:ModerateDemented




Actual:ModerateDemented
Predicted:MildDemented




Actual:MildDemented
Predicted:ModerateDemented




Actual:NonDemented
Predicted:MildDemented



Actual:ModerateDemented
Predicted:ModerateDemented



Actual:MildDemented
Predicted:MildDemented



5. Conclusions

Model performance

DenseNet169 and **VGG16** exhibited the highest performance



Impact and clinical implication

model could serve as a valuable tool for radiologists and neurologists



Future perspectives

1. Refining the deep learning models
2. Data augmentation tasks
3. Collaboration with healthcare professionals and institutions

